



T81-558: Applications of Deep Neural Networks

Module 10: Time Series in Keras

- Instructor: [Jeff Heaton](#), McKelvey School of Engineering, [Washington University in St. Louis](#)
- For more information visit the [class website](#).

Module 10 Material

- Part 10.1: Time Series Data Encoding for Deep Learning [\[Video\]](#) [\[Notebook\]](#)
- Part 10.2: Programming LSTM with Keras and TensorFlow [\[Video\]](#) [\[Notebook\]](#)
- Part 10.3: Text Generation with Keras and TensorFlow [\[Video\]](#) [\[Notebook\]](#)
- Part 10.4: Introduction to Transformers [\[Video\]](#) [\[Notebook\]](#)
- **Part 10.5: Transformers for Timeseries** [\[Video\]](#) [\[Notebook\]](#)

Google CoLab Instructions

The following code ensures that Google CoLab is running the correct version of TensorFlow. Running the following code will map your GDrive to `/content/drive`.

```
In [1]: try:
        from google.colab import drive
        drive.mount('/content/drive', force_remount=True)
        COLAB = True
        print("Note: using Google CoLab")
        %tensorflow_version 2.x
    except:
        print("Note: not using Google CoLab")
        COLAB = False
```

Mounted at /content/drive

Note: using Google CoLab

Part 10.5: Programming Transformers with Keras

This section shows an example of a transformer encoder to predict sunspots. You can find the data files needed for this example at the following location.

- [Sunspot Data Files](#)
- [Download Daily Sunspots](#) - 1/1/1818 to now.

The following code loads the sunspot file:

```
In [2]: import pandas as pd
import os

names = ['year', 'month', 'day', 'dec_year', 'sn_value',
         'sn_error', 'obs_num', 'extra']
df = pd.read_csv(
    "https://data.heatonresearch.com/data/t81-558/SN_d_tot_V2.0.csv",
    sep=';', header=None, names=names,
    na_values=['-1'], index_col=False)

print("Starting file:")
print(df[0:10])

print("Ending file:")
print(df[-10:])
```

Starting file:

| | year | month | day | dec_year | sn_value | sn_error | obs_num | extra |
|---|------|-------|-----|----------|----------|----------|---------|-------|
| 0 | 1818 | 1 | 1 | 1818.001 | -1 | NaN | 0 | 1 |
| 1 | 1818 | 1 | 2 | 1818.004 | -1 | NaN | 0 | 1 |
| 2 | 1818 | 1 | 3 | 1818.007 | -1 | NaN | 0 | 1 |
| 3 | 1818 | 1 | 4 | 1818.010 | -1 | NaN | 0 | 1 |
| 4 | 1818 | 1 | 5 | 1818.012 | -1 | NaN | 0 | 1 |
| 5 | 1818 | 1 | 6 | 1818.015 | -1 | NaN | 0 | 1 |
| 6 | 1818 | 1 | 7 | 1818.018 | -1 | NaN | 0 | 1 |
| 7 | 1818 | 1 | 8 | 1818.021 | 65 | 10.2 | 1 | 1 |
| 8 | 1818 | 1 | 9 | 1818.023 | -1 | NaN | 0 | 1 |
| 9 | 1818 | 1 | 10 | 1818.026 | -1 | NaN | 0 | 1 |

Ending file:

| | year | month | day | dec_year | sn_value | sn_error | obs_num | extra |
|-------|------|-------|-----|----------|----------|----------|---------|-------|
| 72855 | 2017 | 6 | 21 | 2017.470 | 35 | 1.0 | 41 | 0 |
| 72856 | 2017 | 6 | 22 | 2017.473 | 24 | 0.8 | 39 | 0 |
| 72857 | 2017 | 6 | 23 | 2017.475 | 23 | 0.9 | 40 | 0 |
| 72858 | 2017 | 6 | 24 | 2017.478 | 26 | 2.3 | 15 | 0 |
| 72859 | 2017 | 6 | 25 | 2017.481 | 17 | 1.0 | 18 | 0 |
| 72860 | 2017 | 6 | 26 | 2017.484 | 21 | 1.1 | 25 | 0 |
| 72861 | 2017 | 6 | 27 | 2017.486 | 19 | 1.2 | 36 | 0 |
| 72862 | 2017 | 6 | 28 | 2017.489 | 17 | 1.1 | 22 | 0 |
| 72863 | 2017 | 6 | 29 | 2017.492 | 12 | 0.5 | 25 | 0 |
| 72864 | 2017 | 6 | 30 | 2017.495 | 11 | 0.5 | 30 | 0 |

As you can see, there is quite a bit of missing data near the end of the file. We want to find the starting index where the missing data no longer occurs. This technique is somewhat sloppy; it would be better to find a use for the data between missing values.

However, the point of this example is to show how to use a transformer encoder with a somewhat simple time series.

```
In [3]: # Find the last zero and move one beyond
start_id = max(df[df['obs_num'] == 0].index.tolist())+1
print(start_id)
df = df[start_id:] # Trim the rows that have missing observations
```

11314

Divide into training and test/validation sets.

```
In [4]: df['sn_value'] = df['sn_value'].astype(float)
df_train = df[df['year'] < 2000]
df_test = df[df['year'] >= 2000]

spots_train = df_train['sn_value'].tolist()
spots_test = df_test['sn_value'].tolist()

print("Training set has {} observations.".format(len(spots_train)))
print("Test set has {} observations.".format(len(spots_test)))
```

Training set has 55160 observations.

Test set has 6391 observations.

The **to_sequences** function takes linear time series data into an **x** and **y** where **x** is all possible sequences of seq_size. After each **x** sequence, this function places the next value into the **y** variable. These **x** and **y** data can train a time-series neural network.

```
In [5]: import numpy as np

def to_sequences(seq_size, obs):
    x = []
    y = []

    for i in range(len(obs)-SEQUENCE_SIZE):
        #print(i)
        window = obs[i:(i+SEQUENCE_SIZE)]
        after_window = obs[i+SEQUENCE_SIZE]
        window = [[x] for x in window]
        #print("{} - {}".format(window, after_window))
        x.append(window)
        y.append(after_window)

    return np.array(x), np.array(y)

SEQUENCE_SIZE = 10
x_train, y_train = to_sequences(SEQUENCE_SIZE, spots_train)
x_test, y_test = to_sequences(SEQUENCE_SIZE, spots_test)

print("Shape of training set: {}".format(x_train.shape))
print("Shape of test set: {}".format(x_test.shape))
```

Shape of training set: (55150, 10, 1)

Shape of test set: (6381, 10, 1)

We can view the results of the **to_sequences** encoding of the sunspot data.

```
In [6]: print(x_train.shape)
```

```
(55150, 10, 1)
```

Next, we create the transformer_encoder; I obtained this function from a [Keras example](#). This layer includes residual connections, layer normalization, and dropout. This resulting layer can be stacked multiple times. We implement the projection layers with the Keras Conv1D.

```
In [7]: from tensorflow import keras
        from tensorflow.keras import layers

        def transformer_encoder(inputs, head_size, num_heads, ff_dim, dropout=0):
            # Normalization and Attention
            x = layers.LayerNormalization(epsilon=1e-6)(inputs)
            x = layers.MultiHeadAttention(
                key_dim=head_size, num_heads=num_heads, dropout=dropout
            )(x, x)
            x = layers.Dropout(dropout)(x)
            res = x + inputs

            # Feed Forward Part
            x = layers.LayerNormalization(epsilon=1e-6)(res)
            x = layers.Conv1D(filters=ff_dim, kernel_size=1, activation="relu")(x)
            x = layers.Dropout(dropout)(x)
            x = layers.Conv1D(filters=inputs.shape[-1], kernel_size=1)(x)
            return x + res
```

The following function is provided to build the model, including the attention layer.

```
In [8]: def build_model(
        input_shape,
        head_size,
        num_heads,
        ff_dim,
        num_transformer_blocks,
        mlp_units,
        dropout=0,
        mlp_dropout=0,
    ):
        inputs = keras.Input(shape=input_shape)
        x = inputs
        for _ in range(num_transformer_blocks):
            x = transformer_encoder(x, head_size, num_heads, ff_dim, dropout)

        x = layers.GlobalAveragePooling1D(data_format="channels_first")(x)
        for dim in mlp_units:
            x = layers.Dense(dim, activation="relu")(x)
            x = layers.Dropout(mlp_dropout)(x)
```

```
outputs = layers.Dense(1)(x)
return keras.Model(inputs, outputs)
```

We are now ready to build and train the model.

```
In [9]: input_shape = x_train.shape[1:]

model = build_model(
    input_shape,
    head_size=256,
    num_heads=4,
    ff_dim=4,
    num_transformer_blocks=4,
    mlp_units=[128],
    mlp_dropout=0.4,
    dropout=0.25,
)

model.compile(
    loss="mean_squared_error",
    optimizer=keras.optimizers.Adam(learning_rate=1e-4)
)
#model.summary()

callbacks = [keras.callbacks.EarlyStopping(patience=10, \
    restore_best_weights=True)]

model.fit(
    x_train,
    y_train,
    validation_split=0.2,
    epochs=200,
    batch_size=64,
    callbacks=callbacks,
)

model.evaluate(x_test, y_test, verbose=1)
```

```
Epoch 1/200
690/690 [=====] - 25s 16ms/step - loss: 1919.4844 -
val_loss: 463.2157
Epoch 2/200
690/690 [=====] - 11s 16ms/step - loss: 1113.0945 -
val_loss: 365.1375
Epoch 3/200
690/690 [=====] - 11s 16ms/step - loss: 873.6814 -
val_loss: 345.2026
Epoch 4/200
690/690 [=====] - 11s 16ms/step - loss: 789.0035 -
val_loss: 329.4594
Epoch 5/200
690/690 [=====] - 11s 15ms/step - loss: 762.3812 -
val_loss: 324.1521
Epoch 6/200
690/690 [=====] - 11s 15ms/step - loss: 750.8315 -
val_loss: 323.2135
Epoch 7/200
690/690 [=====] - 11s 15ms/step - loss: 744.6664 -
val_loss: 326.0743
Epoch 8/200
690/690 [=====] - 11s 15ms/step - loss: 737.4108 -
val_loss: 312.2708
Epoch 9/200
690/690 [=====] - 11s 15ms/step - loss: 719.6789 -
val_loss: 307.9902
Epoch 10/200
690/690 [=====] - 11s 16ms/step - loss: 715.9462 -
val_loss: 296.0132
Epoch 11/200
690/690 [=====] - 11s 15ms/step - loss: 712.4185 -
val_loss: 299.6908
Epoch 12/200
690/690 [=====] - 11s 15ms/step - loss: 709.5587 -
val_loss: 297.3573
Epoch 13/200
690/690 [=====] - 11s 16ms/step - loss: 703.4562 -
val_loss: 293.3137
Epoch 14/200
690/690 [=====] - 11s 15ms/step - loss: 714.5865 -
val_loss: 310.3690
Epoch 15/200
690/690 [=====] - 11s 15ms/step - loss: 706.6390 -
val_loss: 304.2530
Epoch 16/200
690/690 [=====] - 11s 15ms/step - loss: 701.1292 -
val_loss: 297.4577
Epoch 17/200
690/690 [=====] - 10s 15ms/step - loss: 705.1274 -
val_loss: 326.7051
Epoch 18/200
690/690 [=====] - 11s 15ms/step - loss: 696.4119 -
val_loss: 290.4462
Epoch 19/200
690/690 [=====] - 10s 15ms/step - loss: 701.0396 -
```

```
val_loss: 324.6311
Epoch 20/200
690/690 [=====] - 11s 17ms/step - loss: 694.8000 -
val_loss: 304.3717
Epoch 21/200
690/690 [=====] - 11s 16ms/step - loss: 697.7822 -
val_loss: 314.0597
Epoch 22/200
690/690 [=====] - 11s 15ms/step - loss: 696.5428 -
val_loss: 296.1778
Epoch 23/200
690/690 [=====] - 11s 16ms/step - loss: 688.9620 -
val_loss: 291.3384
Epoch 24/200
690/690 [=====] - 11s 15ms/step - loss: 692.5215 -
val_loss: 294.7356
Epoch 25/200
690/690 [=====] - 10s 15ms/step - loss: 695.5998 -
val_loss: 309.7605
Epoch 26/200
690/690 [=====] - 11s 16ms/step - loss: 688.1234 -
val_loss: 289.6525
Epoch 27/200
690/690 [=====] - 11s 16ms/step - loss: 680.7135 -
val_loss: 287.5633
Epoch 28/200
690/690 [=====] - 11s 15ms/step - loss: 689.2556 -
val_loss: 306.3144
Epoch 29/200
690/690 [=====] - 11s 15ms/step - loss: 686.9375 -
val_loss: 294.5692
Epoch 30/200
690/690 [=====] - 10s 15ms/step - loss: 689.2764 -
val_loss: 295.0640
Epoch 31/200
690/690 [=====] - 11s 15ms/step - loss: 683.3184 -
val_loss: 306.8054
Epoch 32/200
690/690 [=====] - 11s 16ms/step - loss: 679.1677 -
val_loss: 311.3470
Epoch 33/200
690/690 [=====] - 11s 15ms/step - loss: 683.5298 -
val_loss: 292.4295
Epoch 34/200
690/690 [=====] - 11s 15ms/step - loss: 689.2239 -
val_loss: 298.1823
Epoch 35/200
690/690 [=====] - 11s 15ms/step - loss: 682.8902 -
val_loss: 297.4239
Epoch 36/200
690/690 [=====] - 11s 15ms/step - loss: 679.1320 -
val_loss: 289.7046
Epoch 37/200
690/690 [=====] - 11s 16ms/step - loss: 673.3400 -
val_loss: 297.0687
200/200 [=====] - 1s 5ms/step - loss: 214.5603
```

Out[9]: 214.56031799316406

Finally, we evaluate the model with RMSE.

```
In [10]: from sklearn import metrics

pred = model.predict(x_test)
score = np.sqrt(metrics.mean_squared_error(pred,y_test))
print("Score (RMSE): {}".format(score))
```

Score (RMSE): 14.647875946283007